

Operating System Project

An implementation of server-client database using non-blocking operations

Julmy, S., Papinutto, M. & Veillard, S.

University of Fribourg

30 mai 2018

Multi-threaded Server

We selected a multithreaded server as :

- Allows to access same data structure
- light weight as compared to multi-process server
- easier to implement

To use mutex or not to use mutex

Our lock-free implementation allowed us to omit mutex.

- No need to prioritise read or write operations
- No deadlock problems
- All clients have the same right to access the data structure

Atomic calls

Non-blocking operations use atomic operations in order to perform multiple operation in one clock cycle. With GCC, we can access them with some specific compiler functions¹ :

```
type __sync_fetch_and_add(type *ptr, type value, ...);  
bool __sync_bool_compare_and_swap(  
    type* ptr,  
    type old_v,  
    type new_v,  
    ...);
```


Such operations are in a way an acquire lock - operate - release lock in only one CPU cycle.

1. <https://gcc.gnu.org/onlinedocs/gcc-4.1.0/gcc/Atomic-Builtins.html>

Reversed Split-Ordered Hash-Set

This implementation² offers a rapid access to the data but might require slightly more memory than other data structures.

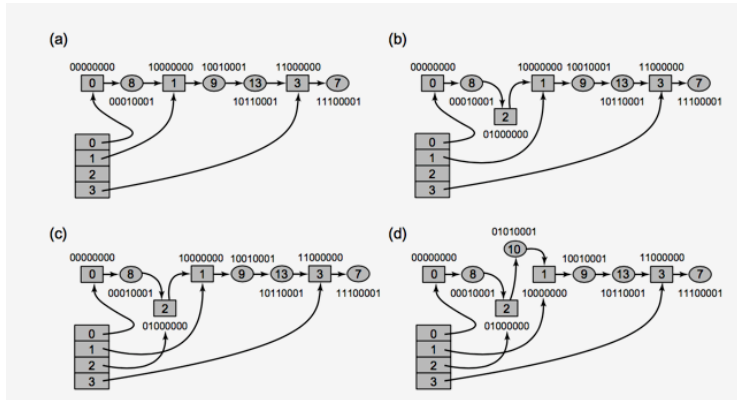
- Buckets are linked in a list which grows automatically when adding elements.
- To expand the set without too much work, we simply add more shortcut in from the first bucket.
- We use sentinel between the nodes to avoid "corner case" that occurs when deleting a reference by a bucket reference.
- The keys are ordered in the reverse binary order : $0x0F101000 \mapsto 0x000808F0$.



2. The Art of multiprocessor programming, Herlihy and Shavit

Operation add in this Hash-Set

Scheme to add the value 10 in the set :



Program Basic Usage

Server usage

TCP Port	5000 (can be changed in file)
<code>./server</code>	server start

Client Start

<code>./ client <server ip address></code>	client start
<code>./ client -option <server ip address></code>	client start with options

Client Usage and Commands

Client Start with options

-?	
-h	
-help	client command help
-f <file>	
-file <file>	client start and execute commands in the file
-F <file1> ... <fileN>	
-files <file1> ... <fileN>	client start and execute commands in the files

Commands in interactive CLI

add <value> or add <key> <value>	add a value to the database
ls	list content (unordered)
read_v <key>	read value from key
read_k <value>	read key from value
rm_v <key>	delete value from key
rm_k <value>	delete value from key
update_kv <value> <newvalue>	update an entry

Demo

DEMO

Tests scenarios

We tested the following scenarios :

- Scenario with collisions : operations that can collide (a client delete a value before another access it).
- Scenario without collision : operations are selected to avoid collisions.
- Scenario with many clients : several clients with a similar scenario as no-collisions.

Collision scenarios

11 clients and 28 commands (308 operations in total)

	Add	Read	Delete
Number of errors	0	22	0
Percentage of errors	0%	7.14%	0%

No-collision scenarios

8 clients and 2700 commands (21600 operations in total)

	Add	Read	Delete
Number of errors	0	0	0
Percentage of errors	0%	0%	0%

Many clients scenarios

32 clients and 300 commands (9600 operations in total)

	Add	Read	Delete
Number of errors	0	0	0
Percentage of errors	0%	0%	0%

This last test required more time than the previous one despite the fact that it has half fewer operations

Encountered problems

- `strtok` vs. `strtok_r`
- lock-free hashset is hard to understand and to implement

Conclusion

- Efficient lock-free implementation
- Understand the lock-free concept and its primitives

Thanks for your attention !